


dqn copy

December 7, 2023

0.1 A2.1a) Run the given online Q-learning algorithm

Go through the implementation in the given Jupyter Notebook. Run online Q-learning, that is, use the `QLearningExperiment` with the `QLearner` class on the `CartPole-v1` environment for 200k steps in the environment.

[]:



0.2 A2.1b) Use a replay buffer in Q-learning

Implement online Q-learning with an experience replay buffer by extending the given skeleton of the `DQNExperiment` class. Train your implementation again in the `CartPole-v1` environment for 200k steps.


```
[ ]: class DQNExperiment(QLearningExperiment):  
    """Experiment that perfoms DQN. You can provide your own learner."""  
  
    def __init__(self, params, model, learner=None, **kwargs):  
        super().__init__(params, model, learner=learner, **kwargs)  
        self.use_last_episode = params.get("use_last_episode", True)  
        self.replay_buffer = TransitionBatch(  
            params.get("replay_buffer_size", int(1e5)),  
            self.runner.transition_format(),  
            batch_size=params.get("batch_size", 1024),  
        )
```

```

def _learn_from_episode(self, episode):
    total_loss = 0
    self.replay_buffer.add(episode["buffer"])
    if len(self.replay_buffer) < self.replay_buffer.batch_size:
        return None
    sample = self.replay_buffer.sample()
    for i in range(self.grad_repeats):
        total_loss += self.learner.train(sample)
    return total_loss / self.grad_repeats

```

```
[ ]:
```



0.3 A2.1c) Implement target networks with hard updates

Extend the QLearning class with target-networks that use a hard update rule. Train your implementation again in the CartPole-v1 environment for 200k steps.

```

[ ]: class QLearnerHardTarget(QLearner):
    def __init__(self, model, params={}):
        super().__init__(model, params)
        self.target_update = params.get("target_update", "hard")
        self.target_update_interval = params.get("target_update_interval", 200)
        self.target_update_calls = 0
        if params.get("target_model", True):
            self.target_model = deepcopy(model)
            for p in self.target_model.parameters():
                p.requires_grad = False
        assert (
            self.target_model is None
            or self.target_update == "soft"
            or self.target_update == "copy"
        ), 'If a target model is specified, it needs to be updated using the_
↪ "soft" or "copy" options.'

    def q_values(self, states, target=False):

```


```

    if target == True and self.target_model is not None:
        return self.target_model(states)
    else:
        return self.model(states)

def target_model_update(self):
    if self.target_model is None:
        pass
    self.target_update_calls += 1
    if self.target_update_calls % self.target_update_interval == 0:
        self.target_model.load_state_dict(self.model.state_dict())

```

[]:



0.4 A2.1d) Implement target networks with soft updates

Extend your implementation with a soft-update rule for the target network and test it in the environment CartPole-v1 for 200k steps.

```

[ ]: class QLearnerSoftTarget(QLearnerHardTarget):
    def __init__(self, model, params={}):
        super().__init__(model, params)
        self.target_update = params.get("target_update", "soft")
        self.soft_target_update_param = params.get("soft_target_update_param", 0.1)

    def target_model_update(self):
        if self.target_model is None:
            pass
        self.target_update_calls += 1
        for t_param, param in zip(
            self.target_model.parameters(), self.model.parameters()
        ):
            t_param.data.copy_(
                (1.0 - self.soft_target_update_param) * t_param.data


```

```

        + self.soft_target_update_param * param.data
    )

```

```
[ ]:
```



0.5 A2.1e) Implement double Q-learning

Extend your implementation with double-Q-learning and test it in the `CartPole-v1` environment for 200k steps.


```
[ ]: class DoubleQLearner(QLearnerSoftTarget):
    def __init__(self, model, params={}):
        super().__init__(model, params)
        self.double_q = params.get("double_q", True)

    def _next_state_values(self, batch):
        with th.no_grad():
            qvalues = self.q_values(batch["next_states"], target=False)
            _, actions = qvalues.max(dim=-1, keepdim=True)
            eval_qvalues = self.q_values(batch["next_states"], target=True)

            return eval_qvalues.gather(dim=-1, index=actions)

```


```
[ ]:
```



0.6 A2.1f) Run double Q-learning on MountainCar

Run your implementation of `DoubleQLearner` on the `MountainCar-v0` environment for $200k$ steps. In all likelihood, your agent will not be able to solve the task (reach the goal), and learning should not pick up at all. Explain why that is.

[]:




The only positive reward in MountainClimber comes from reaching the right top of the hill, which requires a coordinated movement. The chances of a random agent reaching it are very low, so there is no prior experience the model could learn from. A possible solution would be to increase the initial exploration phase, i.e. increase the lower size limit of the replay buffer.

0.7 A2.1g) Run double Q-learning on LunarLander

Run your implementation on the `LunarLander-v2` environment for at least 5 million environment steps (more is better). This can take a while, expect 1-3 hours of computation time. Do you get similar results as shown in the lecture?

[]:



The Kernel crashed while executing code in the the current cell or a previous cell. Please review the code in the cell(s) to identify a possible cause of the failure. Click [here](https://aka.ms/vscodeJupyterKernelCrash) for more info. View Jupyter [log](command:jupyter.viewOutput) for further details.

Because of the memory leak in the Gym environment, the simulation crashed around 120k steps in.